

Project #1 (Due October 11, 2016 – by mid night)

NOTE: This assignment should be done by teams of two students. If you do not have a project-partner, let me know and I will pair you up with one.

One way to increase the clock frequency in the 5-stage pipeline is to split each of the IF, EX and MEM stages into two stages each, and to allow either writing or reading of the register file in one cycle. The resulting pipeline will thus have 8 stages (IF1, IF2, ID, EX1, EX2, MEM1, MEM2, WB) with a structural hazard resulting when an instruction on the WB stage writes into the register file while an instruction in the ID stage reads from the register file. A branch predictor predicts the outcome of a branch in the IF1 stage but the branch condition and the target address are resolved in the EX2 stage. The targets of jump instructions are also resolved in EX2 and are included in the branch target buffer. Forwarding paths (and logic) are provided from the EX2/MEM1 and the MEM2/WB buffers to the ID/EX1 buffer.

Hazards in the 8-stage pipeline can be avoided as follows:

- 1) Structural hazards: if the instruction at WB is trying to write into the register file while the instruction at ID is trying to read from the register file, priority is given to the instruction at WB. The instructions at IF1, IF2 and ID are stalled for one cycle while the instruction at WB is using the register file.
- 2) Data hazards:
 - a. If an instruction in EX1 will write into a register X while the instruction in ID is reading from register X, then the instruction in ID (and subsequent instructions) should stall for one cycle to allow forwarding of the result from EX2/MEM1 to ID/EX1 (in the following cycle). A no-op is injected in EX1/EX2.
 - b. If an instruction in EX2 is a *load* instruction which will write into a register X while the instruction in ID is reading from register X, then the instruction in ID (and subsequent instructions) should stall since it does not have the correct content of register X. A no-op is injected in EX1/EX2.
 - c. If an instruction in MEM1 is a *load* instruction which will write into a register X while the instruction in ID is reading from register X, then the instruction in ID (and subsequent instructions) should stall to allow forwarding of the result from MEM2/WB to ID/EX1 (in the following cycle). A no-op is injected in EX1/EX2.
- 3) Control hazards: If, when a branch is resolved in the EX2 stage, a control hazard is detected, the instructions in the IF1, IF2, ID and EX1 stages are flushed.

The goal of this assignment is to simulate the above 8-stages pipeline and experiment with different pipelined designs using real execution traces made available to you in trace files (file_name.tr). Each trace file is a sequence of trace items, where each trace item represents one instruction executed in the program that has been traced. A trace item is a structure:

```
struct trace_item {
    uint8_t type;           // see below
    uint8_t sReg_a;         // 1st operand
    uint8_t sReg_b;         // 2nd operand
    uint8_t dReg;           // dest. operand
    uint32_t PC;            // program counter
    uint32_t Addr;          // mem. address
};
```

where

```
enum trace_item_type {
    ti_NOP = 0,
    ti_RTYPE,
    ti_ITYPE,
    ti_LOAD,
    ti_STORE,
    ti_BRANCH,
    ti_JTYPE,
    ti_SPECIAL,
    ti_JRTYPE
};
```

The “PC” (program counter) field is the address of the instruction itself. The “type” of an instruction provides the key information about the instruction. Here is a more detailed explanation:

NOP - it's a no-op. No further information is provided.

RTYPE - An R-type instruction.

- sReg_a: first register operand (register name)
- sReg_b: second register operand (register name)
- dReg: destination register name
- PC: program counter of this instruction
- Addr: not used

ITYPE - An I-type instruction that is not LOAD, STORE, or BRANCH.

- sReg_a: first register operand (register name)
- sReg_b: not used
- dReg: destination register name
- PC: program counter of this instruction
- Addr: immediate value

LOAD - a load instruction (memory access)

- sReg_a: first register operand (register name)
- sReg_b: not used
- dReg: destination register name
- PC: program counter of this instruction
- Addr: memory address

STORE - a store instruction (memory access)

- sReg_a: first register operand (register name)
- sReg_b: second register operand (register name)
- dReg: not used
- PC: program counter of this instruction
- Addr: memory address

BRANCH - a branch instruction

- sReg_a: first register operand (register name)
- sReg_b: second register operand (register name)
- dReg: not used
- PC: program counter of this instruction
- Addr: target address

JTYPE - a jump instruction

- sReg_a: not used
- sReg_b: not used
- dReg: not used
- PC: program counter of this instruction
- Addr: target address

SPECIAL - it's a special system call instruction

For now, ignore other fields of this instruction.

JRTYPE - a jump register instruction (used for "return" in functions)

sReg_a: source register (that keeps the target address)
sReg_b: not used
dReg: not used
PC: program counter of this instruction
Addr: target address

To avoid dealing with binary files, you are given a program [CPU.c](#) which reads a trace file (a binary file) and simulates a single cycle CPU (a very simple simulation). It outputs the total number of cycles needed to execute the instructions in the trace file, and if the *trace_view_on* switch is set, outputs also the details of the instruction that finished execution in each cycle. Hence, the program *CPU.c*, which includes [CPU.h](#), takes two arguments; the name of the trace file and a switch value (0 or 1). For example, when you execute “CPU sample.tr 1” you get an output that looks like (if the second argument is 0 rather than 1, only the last line is printed):

```
[cycle 1] LOAD: (PC: 2000a0)(sReg_a: 29)(dReg: 16)(addr: 7fff8000)
[cycle 2] ITYPE: (PC: 2000a4)(sReg_a: 255)(dReg: 28)(addr: 1001)
[cycle 3] ITYPE: (PC: 2000a8)(sReg_a: 28)(dReg: 28)(addr: ffffc000)
[cycle 4] ITYPE: (PC: 2000ac)(sReg_a: 29)(dReg: 17)(addr: 4)
[cycle 5] ITYPE: (PC: 2000b0)(sReg_a: 17)(dReg: 3)(addr: 4)
[cycle 6] ITYPE: (PC: 2000b4)(sReg_a: 255)(dReg: 2)(addr: 2)
[cycle 7] RTYPE: (PC: 2000b8)(sReg_a: 3)(sReg_b: 2)(dReg: 3)
[cycle 8] RTYPE: (PC: 2000bc)(sReg_a: 0)(sReg_b: 3)(dReg: 18)
[cycle 9] STORE: (PC: 2000c0)(sReg_a: 28)(sReg_b: 18)(addr: 10004884)
[cycle 10] ITYPE: (PC: 2000c4)(sReg_a: 29)(dReg: 29)(addr: fffffffe8)
[cycle 11] RTYPE: (PC: 2000c8)(sReg_a: 0)(sReg_b: 16)(dReg: 4)
. . .
[cycle 21] STORE: (PC: 20cd7c)(sReg_a: 29)(sReg_b: 31)(addr: 7fff7fcc)
[cycle 22] BRANCH: (PC: 20cd80)(sReg_a: 16)(sReg_b: 0)(addr: 20cda8)
[cycle 23] LOAD: (PC: 20cd84)(sReg_a: 16)(dReg: 4)(addr: 7fff8007)
. . .
[cycle 32] BRANCH: (PC: 20a9a4)(sReg_a: 17)(sReg_b: 0)(addr: 20a9b4)
[cycle 33] RTYPE: (PC: 20a9b4)(sReg_a: 0)(sReg_b: 0)(dReg: 16)
....

+ Simulation terminates at cycle : 1000
```

The project is to replace the simple single cycle simulation with a simulation of the 8-stages pipeline which will also output the total number of execution cycles as well as the instruction that exits the pipeline in each cycle (if the switch *trace_view_on* is set to 1).

In addition to stalling/flushing due to hazards, your simulation should take a third argument, *prediction_method* (in addition to the trace file name and *trace_view_on*). This argument will be 0, 1, or 2 to reflect three possible designs as follows:

- If *prediction_method* = 0, your simulation should assume that branches are always predicted as “not taken”. In this case, if the prediction is wrong, the four instructions that entered the pipeline before the branch condition is resolved should be flushed.
- If *prediction_method* = 1, your simulation should assume that the architecture uses a one-bit branch predictor which records the last branch condition and address. This predictor is consulted in the IF1 stage which means that if the prediction is correct, the instruction following the branch in the pipeline is the correct one. Otherwise, the wrong prediction will be discovered in the EX2 stage and the four instructions following the branch in the pipeline will be flushed. When no prediction can be made, the “predict not taken” policy should be

assumed. Use a Branch Prediction Hash Table with 64 entries and index this table with bits 9-4 of the branch instruction address (note that some addresses will collide and thus some stored information will be lost – that is OK).

- If *prediction_method* = 2, your simulation should assume that the architecture uses a 2-bit branch predictor.

Your implementation should simulate the 8 pipeline stages and determine the instruction (or an inserted no-op as a result of stalling or flushing) that is in each stage at every cycle. Note that the traces you are given are dynamic traces, hence they do not show which instructions are following a wrong branch (the squashed instructions). You should introduce these squashed instructions into the pipeline without knowing what instructions they were exactly (they were squashed anyways – you should print them as SQUASHED in the output).

TRACES: You are provided with 4 short and 2 long trace files (sample1.tr, sample2.tr, sample3.tr, sample4.tr) and (sample_large1.tr, sample_large2.tr). These files are accessible at [/afs/cs.pitt.edu/courses/1541/long_traces](https://afs.cs.pitt.edu/courses/1541/long_traces) and [/afs/cs.pitt.edu/courses/1541/short_traces](https://afs.cs.pitt.edu/courses/1541/short_traces), where you can also find [sample.tr](#), a small trace file you can use while you debug.

Your assignment is to modify CPU.c to simulate the 8-stages pipeline and test your simulation on the traces provided. You should submit the output of your simulation for each short and long trace file with *trace_view_on* = 0 and *prediction_method* = 0, 1, and 2.

What to submit (email to the TA):

- 1) Your source code for a *CPU.c* simulator which takes 3 arguments; the input trace file, the branch *prediction_method* and the *trace_view_on* switch (in that order). The argument *prediction_method* should be 0 (for “predict not taken”), 1 (for “1-bit branch predictor”), or 2 (for 2-bit predictor). In case the last two parameters are not specified, their default values should be *prediction_method* = 0 and *trace_view_on* = 0.
- 2) The result of running your simulation with *trace_view_on* = 0 for each short and long trace file and *prediction_method* = 0, 1, and 2. Put the results in a table and write down a short analysis of the effect of the branch predictors (observed from the results).
- 3) Change the size of the prediction table (should be a parameter in your program) to 32 and 128 (instead of 64). Compare the results for the three different table sizes and comment on the effect of the size of the prediction table.